

## eda

```
install.packages("GGally")
install.packages("moments")
install.packages("corrplot")
```

```
library(GGally)
library(ggplot2)
library(lmtest)
library(moments)
library(sandwich)
library(stargazer)
library(tidyverse)
library(corrplot)
library(data.table)
library(lubridate)
```

```
# various functions for wrangling
source('./functions/get_robust_se.R')
source('./functions/get_clean_dataset.R')
source('./functions/eda_calculate_stats_by_group.R')
source('./functions/eda_build_quantile_table.R')
```

```
d <- read.csv('data/googleplaystore.csv')
summary(d)
```

```
##      App                Category      Rating      Reviews
## Length:10841      Length:10841   Min.   : 1.000   Length:10841
## Class :character   Class :character 1st Qu.: 4.000   Class :character
## Mode  :character   Mode  :character Median : 4.300   Mode  :character
##                                     Mean  : 4.193
##                                     3rd Qu.: 4.500
##                                     Max.   :19.000
##                                     NA's   :1474
##      Size                Installs      Type      Price
## Length:10841      Length:10841   Length:10841   Length:10841
## Class :character   Class :character   Class :character   Class :character
## Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
##      Content.Rating      Genres      Last.Updated      Current.Ver
## Length:10841      Length:10841   Length:10841   Length:10841
## Class :character   Class :character   Class :character   Class :character
## Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
```

```
##
## Android.Ver
## Length:10841
## Class :character
## Mode :character
##
##
##
##
```

```
d <- get_clean_dataset()
```

```
# summary of dataset
summary(d)
```

```
##      installs      size      reviews      rating
## Min.   :      500  Min.   : 0.85  Min.   :   100  Min.   :1.600
## 1st Qu.: 100000  1st Qu.: 7.10  1st Qu.:   929  1st Qu.:4.000
## Median : 1000000  Median : 19.00  Median :   9194  Median :4.300
## Mean   : 10215340  Mean   : 26.89  Mean   : 373798  Mean   :4.198
## 3rd Qu.: 5000000  3rd Qu.: 40.00  3rd Qu.:  71762  3rd Qu.:4.500
## Max.   :1000000000  Max.   :100.00  Max.   :44893888  Max.   :5.000
##      price      is_free      last_updated      android_version
## Min.   : 0.000  Mode :logical  Min.   :0.00000  Min.   :1.000
## 1st Qu.: 0.000  FALSE:305     1st Qu.:0.04931  1st Qu.:4.000
## Median : 0.000  TRUE :5103    Median :0.19178  Median :4.100
## Mean   : 1.036                      Mean   :0.71197  Mean   :3.845
## 3rd Qu.: 0.000                      3rd Qu.:0.83836  3rd Qu.:4.100
## Max.   :400.000                      Max.   :8.21644  Max.   :8.000
## current_version  category      is_family_category  is_game_category
## Min.   : 0.000  Length:5408    Mode :logical      Mode :logical
## 1st Qu.: 1.200  Class :character  FALSE:4245         FALSE:4554
## Median : 2.100  Mode :character  TRUE :1163         TRUE :854
## Mean   : 5.157
## 3rd Qu.: 4.100
## Max.   :858.000
## is_tools_category  genre      content_rating      is_content_everyone
## Mode :logical      Length:5408    Length:5408         Mode :logical
## FALSE:4996         Class :character  Class :character    FALSE:1228
## TRUE :412          Mode :character  Mode :character     TRUE :4180
##
##
##
##      type      install_group      log_installs      log_size
## Length:5408    Min.   : 1.00  Min.   :2.699  Min.   : -0.07058
## Class :character  1st Qu.: 6.00  1st Qu.:5.000  1st Qu.: 0.85126
## Mode :character  Median : 8.00  Median :6.000  Median : 1.27875
##                      Mean   : 7.29  Mean   :5.704  Mean   : 1.21414
##                      3rd Qu.: 9.00  3rd Qu.:6.699  3rd Qu.: 1.60206
##                      Max.   :14.00  Max.   :9.000  Max.   : 2.00000
## log_current_version  log_last_updated  log_reviews
## Min.   :0.0000  Min.   :0.00000  Min.   :2.000
## 1st Qu.:0.3424  1st Qu.:0.02091  1st Qu.:2.968
## Median :0.4914  Median :0.07620  Median :3.964
## Mean   :0.5673  Mean   :0.17347  Mean   :3.995
```

```
## 3rd Qu.:0.7076      3rd Qu.:0.26443      3rd Qu.:4.856
## Max.      :2.9340      Max.      :0.96456      Max.      :7.652

# save a data.table version for some easier wrangling downstream
d_dt <- as.data.table(d)

numeric_cols <- colnames(d)[unlist(lapply(d, is.numeric))]

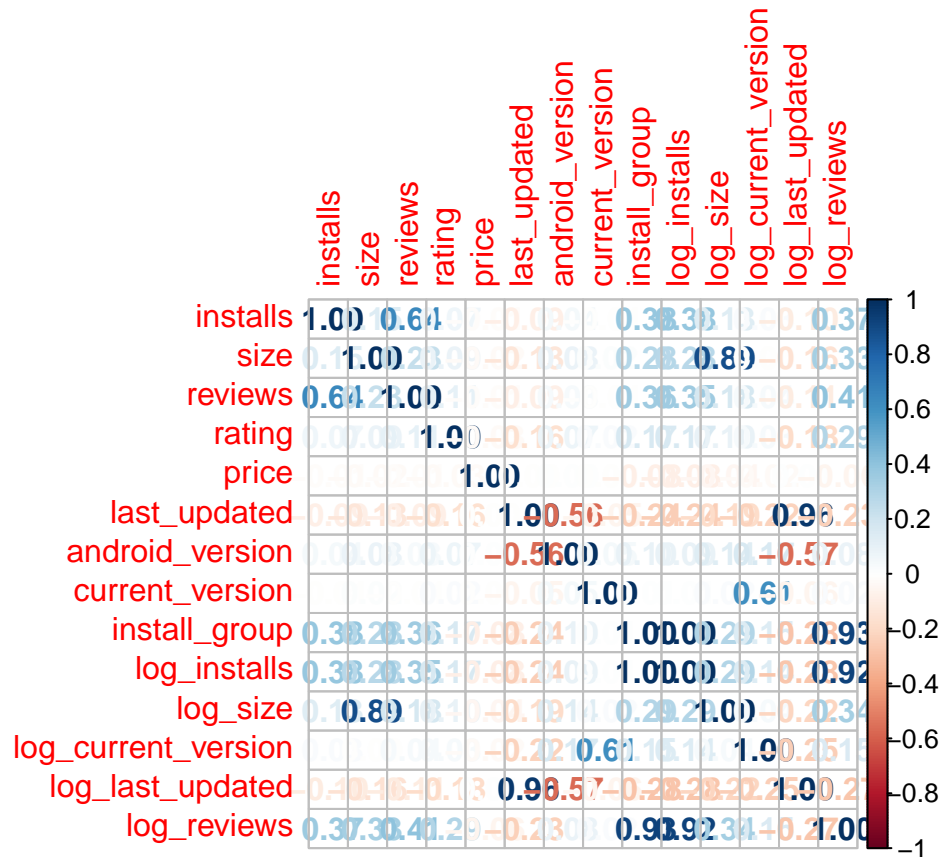
table_quantile_numeric <- rbindlist(lapply(numeric_cols, eda_build_quantile_table))
table_quantile_numeric
```

##	0%	5%	25%	50%	
## 1:	500.00000000	10000.00000000	100000.00000000	1000000.00000000	
## 2:	0.85000000	2.50000000	7.10000000	19.00000000	
## 3:	100.00000000	156.00000000	928.75000000	9194.00000000	
## 4:	1.60000000	3.40000000	4.00000000	4.30000000	
## 5:	0.00000000	0.00000000	0.00000000	0.00000000	
## 6:	0.00000000	0.01369863	0.04931507	0.19178082	
## 7:	1.00000000	2.20000000	4.00000000	4.10000000	
## 8:	0.00000000	1.00000000	1.20000000	2.10000000	
## 9:	1.00000000	4.00000000	6.00000000	8.00000000	
## 10:	2.69897000	4.00000000	5.00000000	6.00000000	
## 11:	-0.07058107	0.39794001	0.85125835	1.27875360	
## 12:	0.00000000	0.30103000	0.34242268	0.49136169	
## 13:	0.00000000	0.00590886	0.02090591	0.07619639	
## 14:	2.00000000	2.19312460	2.96789878	3.96350444	
##	75%	95%	100%		variable
## 1:	5000000.00000000	50000000.00000000	1000000000.00000000		installs
## 2:	40.00000000	82.00000000	100.00000000		size
## 3:	71762.25000000	1224897.70000000	44893888.00000000		reviews
## 4:	4.50000000	4.70000000	5.00000000		rating
## 5:	0.00000000	0.99000000	400.00000000		price
## 6:	0.8383562	3.2884932	8.2164384		last_updated
## 7:	4.10000000	5.00000000	8.00000000		android_version
## 8:	4.10000000	10.00000000	858.00000000		current_version
## 9:	9.00000000	11.00000000	14.00000000		install_group
## 10:	6.6989700	7.6989700	9.00000000		log_installs
## 11:	1.6020600	1.9138139	2.00000000		log_size
## 12:	0.7075702	1.0413927	2.9339932		log_current_version
## 13:	0.2644297	0.6323046	0.9645631		log_last_updated
## 14:	4.8558960	6.0880991	7.6521872		log_reviews
##	diff_min_vs_max				
## 1:	999999500.0000000				
## 2:	99.1500000				
## 3:	44893788.0000000				
## 4:	3.4000000				
## 5:	400.0000000				
## 6:	8.2164384				
## 7:	7.0000000				
## 8:	858.0000000				
## 9:	13.0000000				
## 10:	6.3010300				
## 11:	2.0705811				
## 12:	2.9339932				
## 13:	0.9645631				

```
## 14: 5.6521872
```

```
# Corrplot across variables
corrplot(cor(d[,numeric_cols], use = "complete.obs"),
         method = 'number')


# Corrplot across variables (5th PCTL outliers removed)
corrplot(cor(d[d$reviews >= 6,numeric_cols], use = "complete.obs"),
         method = 'number')
```



```
# Corrplot across variables (25th PCTL outliers removed)
corrplot(cor(d[d$reviews >= 100,numeric_cols], use = "complete.obs"),
         method = 'number')


# Distribution of numeric columns
ggplot(data = d, aes(x = rating)) +
  geom_histogram()

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = log(rating))) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = installs)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
ggplot(data = d, aes(x = log(installs))) +  
  geom_histogram()
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
ggplot(data = d, aes(x = install_group)) +  
  geom_histogram()
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```






```
ggplot(data = d, aes(x = log(install_group))) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = reviews)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = log(reviews))) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = size)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = log(size))) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
ggplot(data = d, aes(x = price)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
ggplot(data = d, aes(x = log(price))) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
## Warning: Removed 5103 rows containing non-finite values (stat_bin).
```



```
categorical_cols <- c(
  'category',
  'type',
  'content_rating',
  'current_version',
  'android_version'
)
```

```
# perform function across all categorical columns
```

```
table_long_cat <- rbindlist(lapply(categorical_cols, eda_calculate_stats_by_group))
```

```
table_quantile_cat <- rbindlist(lapply(categorical_cols, eda_calculate_stats_by_group, quantile_table=TRUE))
```

```
# compare mean install grp across variable values
```

```
table_long_cat
```

##	group_by_column	count_apps	install_group_avg	install_count_med
## 1:	GAME	854	8.553864	5000000
## 2:	PHOTOGRAPHY	182	8.456044	1000000
## 3:	SHOPPING	132	8.196970	1000000
## 4:	COMMUNICATION	144	7.972222	1000000
## 5:	PRODUCTIVITY	156	7.474359	1000000
## 6:	SOCIAL	123	7.398374	1000000
## 7:	SPORTS	195	7.158974	1000000
## 8:	HEALTH_AND_FITNESS	150	7.140000	500000
## 9:	NEWS_AND_MAGAZINES	108	7.083333	1000000
## 10:	PERSONALIZATION	184	7.027174	750000



## 11:	FAMILY	1163	7.015477	500000
## 12:	TOOLS	412	7.002427	500000
## 13:	LIFESTYLE	171	6.520468	100000
## 14:	DATING	108	6.518519	500000
## 15:	BUSINESS	114	6.447368	100000
## 16:	FINANCE	186	6.107527	100000
## 17:	MEDICAL	125	5.408000	100000
## 18:	Free	5103	7.461885	1000000
## 19:	Paid	305	4.419672	10000
## 20:	Everyone 10+	269	8.189591	1000000
## 21:	Teen	685	7.854015	1000000
## 22:	Mature 17+	271	7.236162	500000
## 23:	Everyone	4180	7.144019	500000
## 24:	1.4	175	7.708571	1000000
## 25:	3.1	131	7.610687	1000000
## 26:	1.7	107	7.560748	1000000
## 27:	2.2	138	7.557971	1000000
## 28:	1.5	129	7.054264	500000
## 29:	1.2	320	7.046875	500000
## 30:	1.1	403	7.019851	500000
## 31:	1.6	106	6.962264	500000
## 32:	2.1	207	6.888889	500000
## 33:	1.3	207	6.879227	500000
## 34:	3	128	6.804688	100000
## 35:	2	211	6.625592	500000
## 36:	1	641	6.262090	100000
## 37:	4.1	1417	7.783345	1000000
## 38:	4	1616	7.321782	1000000
## 39:	4.4	572	7.295455	1000000
## 40:	4.2	230	7.265217	1000000
## 41:	5	324	7.166667	1000000
## 42:	4.3	126	7.039683	500000
## 43:	2.3	611	6.978723	500000
## 44:	3	137	6.583942	100000
## 45:	2.2	124	5.750000	100000
##	group_by_column	count_apps	install_group_avg	install_count_med
##	variable			
## 1:	category			
## 2:	category			
## 3:	category			
## 4:	category			
## 5:	category			
## 6:	category			
## 7:	category			
## 8:	category			
## 9:	category			
## 10:	category			
## 11:	category			
## 12:	category			
## 13:	category			
## 14:	category			
## 15:	category			
## 16:	category			
## 17:	category			

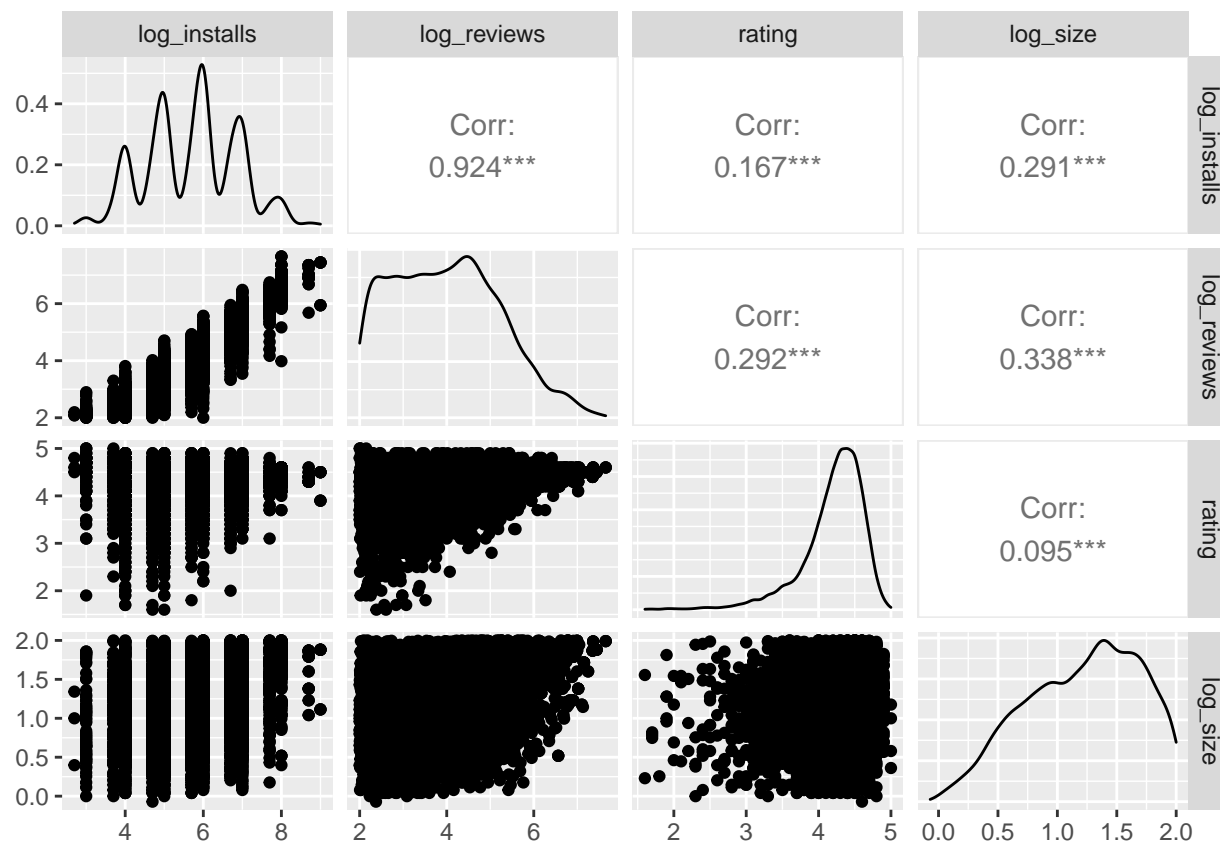
```
## 18:          type
## 19:          type
## 20: content_rating
## 21: content_rating
## 22: content_rating
## 23: content_rating
## 24: current_version
## 25: current_version
## 26: current_version
## 27: current_version
## 28: current_version
## 29: current_version
## 30: current_version
## 31: current_version
## 32: current_version
## 33: current_version
## 34: current_version
## 35: current_version
## 36: current_version
## 37: android_version
## 38: android_version
## 39: android_version
## 40: android_version
## 41: android_version
## 42: android_version
## 43: android_version
## 44: android_version
## 45: android_version
##          variable
```

```
# compare distribution of mean install grp across variable values
table_quantile_cat
```

##	0%	25%	50%	75%	100%	variable	diff_min_vs_max
## 1:	5.408000	6.520468	7.083333	7.474359	8.553864	category	3.145864
## 2:	4.419672	5.180225	5.940779	6.701332	7.461885	type	3.042213
## 3:	7.144019	7.213127	7.545088	7.937909	8.189591	content_rating	1.045572
## 4:	6.262090	6.879227	7.019851	7.557971	7.708571	current_version	1.446481
## 5:	5.750000	6.978723	7.166667	7.295455	7.783345	android_version	2.033345

**I.I.D. data:** According to the Kaggle authors, this data set was collected by randomly scraping the Google Play Store. Since no clusters of applications were specifically targeted, we can reasonably use the entirety of the store as our reference population. We recognize that applications likely have some degree of interdependence, especially within genres. For example, the success of one application probably has a negative impact on other applications of the same type. Due to the large size of this data set (5408 records), however, we expect any dependencies to be negligible. We also have reason to believe that the data are identically distributed, as they are drawn from the same population of applications. One could argue that since the Google Play Store changes over time, the distribution also shifts in response. Because the authors do not mention the time frame across which the data was collected, we will assume that they originated from a single snapshot of the Play Store and that no shifts in the underlying distribution occurred.

```
cols <- c('log_installs', 'log_reviews', 'rating', 'log_size')
ggpairs(d[, cols])
```



```
# save a data.table version for some easier wrangling downstream
```

```
d_dt <- as.data.table(d)
```


```
cols <- c('rating', 'install_group', 'installs', 'log_installs', 'log_reviews', 'log_size')
```

```
# Corrplot across variables
```

```
corrplot(cor(d[,  
             cols], use = "complete.obs"),  
         method = 'number')
```

```
# Corrplot across variables (5th PCTL outliers removed)
```

```
corrplot(cor(d[d$reviews >= 6, cols], use = "complete.obs"),  
         method = 'number')
```



```
# Corrplot across variables (25th PCTL outliers removed)
corrplot(cor(d[d$reviews >= 100,cols], use = "complete.obs"),
          method = 'number')
```

```
model_small <- lm(log_installs ~ 1 + log_reviews, data = d)
model_medium <- lm(log_installs ~ 1 + log_reviews + rating + log_size, data = d)
model_large <- lm(log_installs ~ 1 + log_reviews + rating + log_size + factor(type), data = d)

stargazer(
  model_small,
  model_medium,
  model_large,
  type = 'text',
  se = list(get_robust_se(model_small), get_robust_se(model_medium))
)
```

```
##
## =====
##                               Dependent variable:
##                               -----
##                               log_installs
##                               (1)           (2)           (3)
## -----
## log_reviews                0.879***      0.918***      0.889***
##                           (0.005)      (0.005)      (0.005)
##
## rating                     -0.302***      -0.253***
```


```
##                                     (0.017)                      (0.013)
##
## log_size                          -0.061***                  -0.049***
##                                     (0.013)                      (0.012)
##
## factor(type)Paid                  -0.720***
##                                     (0.024)
##
## Constant                          2.194***                  3.379***
##                                     (0.021)                  (0.069)
##                                     (0.053)
## -----
## Observations                      5,408                      5,408
## R2                                0.854                      0.866
## Adjusted R2                       0.854                      0.866
## Residual Std. Error    0.443 (df = 5406)    0.425 (df = 5404)    0.392 (df = 5403)
## F Statistic           31,576.970*** (df = 1; 5406) 11,635.130*** (df = 3; 5404) 10,452.460*** (df = 4;
## =====
## Note:                                     *p<0.1; **p<0.05; ***p<0.01
```

2. **No Perfect Colinearity:** We can immediately conclude that `log_installs`, `log_reviews`, `rating`, and `log_size` are not perfectly colinear as otherwise the regression above would have failed. We can also assess near perfect colinearity for these variables by observing the robust standard errors returned by the regression model. In general, highly colinear features will have large standard errors. Since the standard error of the coefficients are small relative to their magnitude, we can reasonably conclude that they are not nearly colinear.

3. **Linear Conditional Expectation:** To verify the assumption of linear conditional expectations, we seek to show that there is no relationship between the model residuals and any of the predictor variables. That is, the model does not systematically underpredict or overpredict in certain regions of the input space. Plots 1 through 3 show the relationships between the model residuals and individual predictors. The residuals are generally well-centered around zero, although the model seems to underpredict when `log_reviews` is high and `rating` is low. The fourth plot shows the model residuals as a function of the model predictions. Here, the model seems to underpredict in the left-most and right-most regions, and slightly overpredict in the middle. Overall, there are no strong non-linear relationships between the model residuals and the input features, and we do not find enough evidence to reject the assumption of linear conditional expectation.


```
# Reviews versus residuals
plot_1 <- ggplot(data = d, mapping = aes(x = log_reviews, y = resid(model_medium))) +
  geom_point() + stat_smooth()
plot_1

## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```




```
# Ratings versus residuals  
plot_2 <- ggplot(data = d, mapping = aes(x = rating, y = resid(model_medium))) +  
  geom_point() + stat_smooth()  
plot_2
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



```
# Size versus residuals
plot_3 <- ggplot(data = d, mapping = aes(x = log_size, y = resid(model_medium))) +
  geom_point() + stat_smooth()
plot_3
```


```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



```
# Model predictions versus residuals  
plot_4 <- ggplot(data = d, mapping = aes(x = predict(model_medium), y = resid(model_medium))) +  
  geom_point() + stat_smooth()  
plot_4
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```






4. **Homoskedastic Errors:** When assessing homoskedastic errors, we seek to determine if there is a relationship between the variance of the model residuals and the predictors. If the homoskedastic assumption is satisfied, then we should observe a lack of relationship; conversely, if the data are heteroskedastic then the conditional variance will depend on the predictors. The first plot is an eyeball test of homoskedasticity, showing the model residuals as a function of the model predictions. We notice that the spread of the residuals is mostly consistent throughout the data, although the right-hand side is somewhat narrower. As a more concrete assessment, we also perform a Breusch-Pagan test with the null hypothesis that there are no heteroskedastic errors in the model. Since the  $p$ -value falls below our significance threshold of 0.001, we find enough evidence to reject the null hypothesis. In response to this failed assumption, we report robust standard errors (adjusted for heteroskedasticity) instead of non-adjusted errors.

```
# Breusch-Pagan test
bp_test <- bptest(model_small)
bp_test

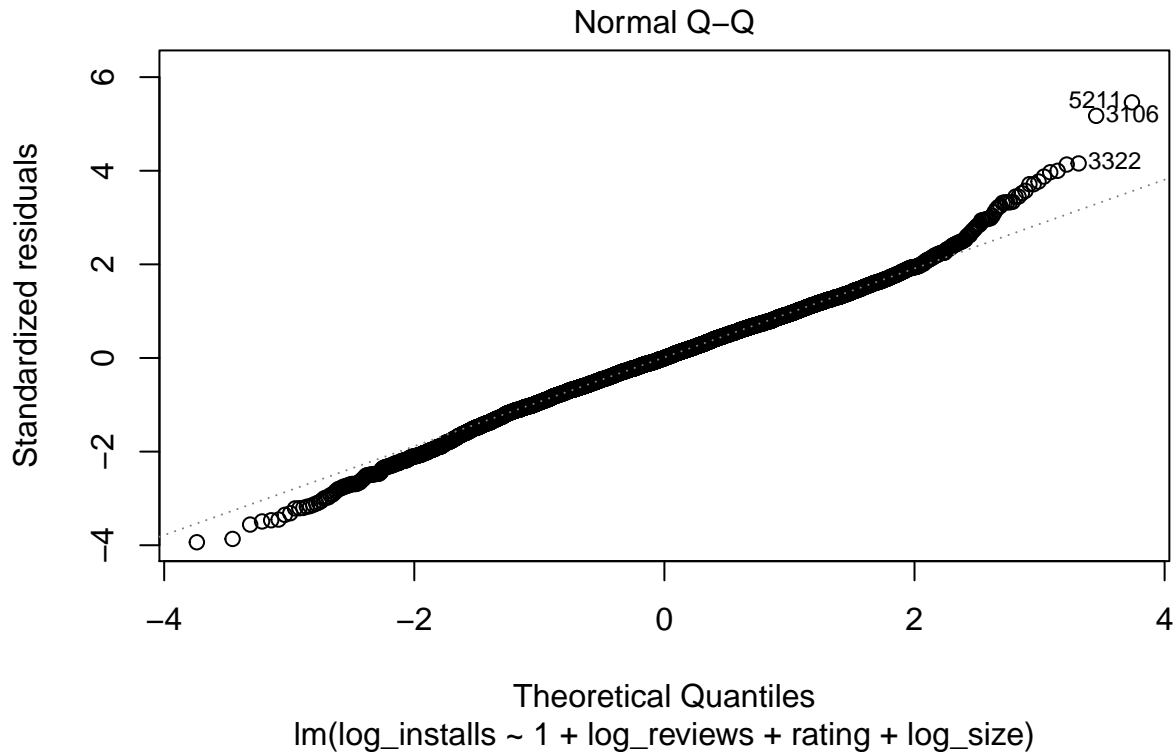
##
## studentized Breusch-Pagan test
##
## data: model_small
## BP = 107.58, df = 1, p-value < 0.00000000000000022
plot_4

## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



5. **Normally Distributed Errors:** When assessing the normality of the error distribution, we seek to determine if the model residuals are approximately Gaussian. If so, then the sample quantiles of the residuals should closely match the theoretical quantiles of a normal distribution in a Q-Q plot. Below, we plot the Q-Q plot associated with our model. In general, the residuals seem to follow a normal distribution, as the middle quantiles match the corresponding theoretical quantiles. However, the tails of the residual distribution are fatter than expected; the first quantiles occur at smaller than expected values, and the last quantiles occur at larger than expected values. Overall, the assumption of normally distributed errors seems imperfect but reasonably justified.

```
# Q-Q plot
plot_5 <- plot(model_medium, which = 2)
```



```
plot_5
```

```
## NULL
```

**\*\* Reverse Causality: \*\*** We have to consider the possibility that high average reviews could lead to a higher number of installations which could lead to a higher average review. We will want to test for a reverse causality relationship between these two variables to determine if the best linear predictor is valid. If we regress average reviews on installs, the installs coefficient ( $\gamma_1$ ) will have a positive slope. Since  $\beta_1$  (average review slope coefficient)  $> 0$ , we know higher average review leads to more installs. Since  $\gamma_1$  (installs slope coefficient for reverse causality) is  $> 0$ , this leads to positive feedback. Given we have two potentially positive coefficients, this could be a bias away from zero which is a concern that a reverse causality relationship exists between the two variables. We could consider dropping average reviews as a variable and determine if there are other leading variables that can explain the number of installs for an app.

```
model_small <- lm(log_installs ~ 1 + log_reviews, data = d)
model_reverse <- lm(log_reviews ~ 1 + log_installs, data = d)

stargazer(
  model_small,
  model_reverse,
  type = 'text',
  se = list(get_robust_se(model_small), get_robust_se(model_medium))
)
```

```
##
## =====
##                               Dependent variable:
```

```

##          -----
##          log_installs   log_reviews
##          (1)           (2)
## -----
## log_reviews          0.879***
##                   (0.005)
##
## log_installs                0.972
##
##
## Constant          2.194***   -1.548***
##                   (0.021)   (0.069)
##
## -----
## Observations          5,408       5,408
## R2                    0.854       0.854
## Adjusted R2           0.854       0.854
## Residual Std. Error (df = 5406)  0.443       0.466
## F Statistic (df = 1; 5406)  31,576.970***  31,576.970***
## =====
## Note:                *p<0.1; **p<0.05; ***p<0.01

```